



Transformers documentation

Token classification ▾

Token classification



Open in Colab



Open

Studio Lab

Copy page



Tasks: Token Classification



Token classification assigns a label to individual tokens in a sentence. One of the most common token classification tasks is Named Entity Recognition (NER). NER attempts to find a label for each entity in a sentence, such as a person, location, or organization.

This guide will show you how to:

1. Finetune [DistilBERT](#) on the [WNUT 17](#) dataset to detect new entities.
2. Use your finetuned model for inference.

To see all architectures and checkpoints compatible with this task, we recommend checking the [task page](#).

Before you begin, make sure you have all the necessary libraries installed:

We encourage you to login to your Hugging Face account so you can upload and share your model with the community. When prompted, enter your token to login:

Load WNUT 17 dataset

2/11

```
"B-location",  
"I-location",  
"B-person",  
"I-person",  
"B-product",  
"I-product",  
]
```

The letter that prefixes each `ner_tag` indicates the token position of the entity:

- B- indicates the beginning of an entity.
- I- indicates a token is contained inside the same entity (for example, the `State` token is a part of an entity like `Empire State Building`).
- `0` indicates the token doesn't correspond to any entity.

Preprocess

Data processing for Token Classification



The next step is to load a DistilBERT tokenizer to preprocess the `tokens` field:

```
>>> from transformers import AutoTokenizer  
  
>>> tokenizer = AutoTokenizer.from_pretrained("distilbert/distilbert-base-uncased")
```

As you saw in the example `tokens` field above, it looks like the input has already been tokenized. But the input actually hasn't been tokenized yet and you'll need to set `is_split_into_words=True` to tokenize the words into subwords. For example:

```
>>> example = wnut["train"][0]
>>> tokenized_input = tokenizer(example["tokens"], is_split_into_words=True)
>>> tokens = tokenizer.convert_ids_to_tokens(tokenized_input["input_ids"])
>>> tokens
['[CLS]', '@', 'paul', '##walk', 'it', '', 's', 'the', 'view', 'from', 'where', 'i',
```

However, this adds some special tokens `[CLS]` and `[SEP]` and the subword tokenization creates a mismatch between the input and labels. A single word corresponding to a single label may now be split into two subwords. You'll need to realign the tokens and labels by:

1. Mapping all tokens to their corresponding word with the `word_ids` method.
2. Assigning the label `-100` to the special tokens `[CLS]` and `[SEP]` so they're ignored by the PyTorch loss function (see [CrossEntropyLoss](#)).
3. Only labeling the first token of a given word. Assign `-100` to other subtokens from the same word.

Here is how you can create a function to realign the tokens and labels, and truncate sequences to be no longer than DistilBERT's maximum input length:

```
>>> def tokenize_and_align_labels(examples):
...     tokenized_inputs = tokenizer(examples["tokens"], truncation=True, is_split_into_words=True)

...     labels = []
...     for i, label in enumerate(examples[f"ner_tags"]):
...         word_ids = tokenized_inputs.word_ids(batch_index=i)  # Map tokens to their
...         previous_word_idx = None
...         label_ids = []
...         for word_idx in word_ids:  # Set the special tokens to -100.
...             if word_idx is None:
...                 label_ids.append(-100)
...             elif word_idx != previous_word_idx:  # Only label the first token of a
...                 label_ids.append(label[word_idx])
...             else:
...                 label_ids.append(-100)
...             previous_word_idx = word_idx
...         labels.append(label_ids)
```

```
...     tokenized_inputs["labels"] = labels
...     return tokenized_inputs
```

To apply the preprocessing function over the entire dataset, use 🧐 Datasets map function. You can speed up the `map` function by setting `batched=True` to process multiple elements of the dataset at once:

```
>>> tokenized_wnut = wnud.map(tokenize_and_align_labels, batched=True)
```

Now create a batch of examples using DataCollatorWithPadding. It's more efficient to *dynamically pad* the sentences to the longest length in a batch during collation, instead of padding the whole dataset to the maximum length.

```
>>> from transformers import DataCollatorForTokenClassification

>>> data_collator = DataCollatorForTokenClassification(tokenizer=tokenizer)
```

Evaluate

Including a metric during training is often helpful for evaluating your model's performance. You can quickly load a evaluation method with the 🧐 Evaluate library. For this task, load the segeval framework (see the 🧐 Evaluate quick tour to learn more about how to load and compute a metric). Segeval actually produces several scores: precision, recall, F1, and accuracy.

```
>>> import evaluate

>>> segeval = evaluate.load("segeval")
```

Get the NER labels first, and then create a function that passes your true predictions and true labels to compute to calculate the scores:

```
>>> import numpy as np

>>> labels = [label_list[i] for i in example[f"ner_tags"]]
```

```
>>> def compute_metrics(p):
...     predictions, labels = p
...     predictions = np.argmax(predictions, axis=2)

...     true_predictions = [
...         [label_list[p] for (p, l) in zip(prediction, label) if l != -100]
...         for prediction, label in zip(predictions, labels)
...     ]
...     true_labels = [
...         [label_list[l] for (p, l) in zip(prediction, label) if l != -100]
...         for prediction, label in zip(predictions, labels)
...     ]

...     results = seqeval.compute(predictions=true_predictions, references=true_labels)
...     return {
...         "precision": results["overall_precision"],
...         "recall": results["overall_recall"],
...         "f1": results["overall_f1"],
...         "accuracy": results["overall_accuracy"],
...     }
```

Your `compute_metrics` function is ready to go now, and you'll return to it when you setup your training.

Train

Before you start training your model, create a map of the expected ids to their labels with `id2label` and `label2id`:

```
>>> id2label = {
...     0: "O",
...     1: "B-corporation",
...     2: "I-corporation",
...     3: "B-creative-work",
...     4: "I-creative-work",
...     5: "B-group",
...     6: "I-group",
...     7: "B-location",
...     8: "I-location",
...     9: "B-person",
...     10: "I-person",
...     11: "B-product",
...     12: "I-product",
```

```
... }  
>>> label2id = {  
...     "0": 0,  
...     "B-corporation": 1,  
...     "I-corporation": 2,  
...     "B-creative-work": 3,  
...     "I-creative-work": 4,  
...     "B-group": 5,  
...     "I-group": 6,  
...     "B-location": 7,  
...     "I-location": 8,  
...     "B-person": 9,  
...     "I-person": 10,  
...     "B-product": 11,  
...     "I-product": 12,  
... }
```

If you aren't familiar with finetuning a model with the [Trainer](#), take a look at the basic tutorial [here](#)!

You're ready to start training your model now! Load DistilBERT with [AutoModelForTokenClassification](#) along with the number of expected labels, and the label mappings:

```
>>> from transformers import AutoModelForTokenClassification, TrainingArguments, Train  
  
>>> model = AutoModelForTokenClassification.from_pretrained(  
...     "distilbert/distilbert-base-uncased", num_labels=13, id2label=id2label, label2  
... )
```

At this point, only three steps remain:

1. Define your training hyperparameters in [TrainingArguments](#). The only required parameter is `output_dir` which specifies where to save your model. You'll push this model to the Hub by setting `push_to_hub=True` (you need to be signed in to Hugging Face to upload your model). At the end of each epoch, the [Trainer](#) will evaluate the seqeval scores and save the training checkpoint.
2. Pass the training arguments to [Trainer](#) along with the model, dataset, tokenizer, data collator, and `compute_metrics` function.

3. Call `train()` to finetune your model.

```
>>> training_args = TrainingArguments(  
...     output_dir="my_awesome_wnut_model",  
...     learning_rate=2e-5,  
...     per_device_train_batch_size=16,  
...     per_device_eval_batch_size=16,  
...     num_train_epochs=2,  
...     weight_decay=0.01,  
...     eval_strategy="epoch",  
...     save_strategy="epoch",  
...     load_best_model_at_end=True,  
...     push_to_hub=True,  
... )  
  
>>> trainer = Trainer(  
...     model=model,  
...     args=training_args,  
...     train_dataset=tokenized_wnut["train"],  
...     eval_dataset=tokenized_wnut["test"],  
...     processing_class=tokenizer,  
...     data_collator=data_collator,  
...     compute_metrics=compute_metrics,  
... )  
  
>>> trainer.train()
```

Once training is completed, share your model to the Hub with the `push_to_hub()` method so everyone can use your model:

```
>>> trainer.push_to_hub()
```

For a more in-depth example of how to finetune a model for token classification, take a look at the corresponding [PyTorch notebook](https://huggingface.co/docs/transformers/tasks/token_classification).

Inference

Great, now that you've finetuned a model, you can use it for inference!

Grab some text you'd like to run inference on:


```
>>> text = "The Golden State Warriors are an American professional basketball team bas"
```

The simplest way to try out your finetuned model for inference is to use it in a `pipeline()`.

Instantiate a pipeline for NER with your model, and pass your text to it:

```
>>> from transformers import pipeline

>>> classifier = pipeline("ner", model="stevhliu/my_awesome_wnut_model")
>>> classifier(text)
[{'entity': 'B-location',
  'score': 0.42658573,
  'index': 2,
  'word': 'golden',
  'start': 4,
  'end': 10},
 {'entity': 'I-location',
  'score': 0.35856336,
  'index': 3,
  'word': 'state',
  'start': 11,
  'end': 16},
 {'entity': 'B-group',
  'score': 0.3064001,
  'index': 4,
  'word': 'warriors',
  'start': 17,
  'end': 25},
 {'entity': 'B-location',
  'score': 0.65523505,
  'index': 13,
  'word': 'san',
  'start': 80,
  'end': 83},
 {'entity': 'B-location',
  'score': 0.4668663,
  'index': 14,
  'word': 'francisco',
  'start': 84,
  'end': 93}]
```

You can also manually replicate the results of the `pipeline` if you'd like:

Tokenize the text and return PyTorch tensors:

```
>>> from transformers import AutoTokenizer

>>> tokenizer = AutoTokenizer.from_pretrained("stevhliu/my_awesome_wnut_model")
>>> inputs = tokenizer(text, return_tensors="pt")
```

Pass your inputs to the model and return the logits:

```
>>> from transformers import AutoModelForTokenClassification

>>> model = AutoModelForTokenClassification.from_pretrained("stevhliu/my_awesome_wnut_
>>> with torch.no_grad():
...     logits = model(**inputs).logits
```

Get the class with the highest probability, and use the model's `id2label` mapping to convert it to a text label:

```
>>> predictions = torch.argmax(logits, dim=2)
>>> predicted_token_class = [model.config.id2label[t.item()]] for t in predictions[0]]
>>> predicted_token_class
['0',
 '0',
 'B-location',
 'I-location',
 'B-group',
 '0',
 '0',
 '0',
 '0',
 '0',
 '0',
 '0',
 '0',
 '0',
 '0',
 'B-location',
 'B-location',
 '0',
 '0']
```

[Update on GitHub](#)

← Text classification

Question answering →